

Causal DAG extraction from a library of books or videos/movies

Robert R. Tucci
tucci@ar-tiste.com

October 29, 2022

Abstract

Determining a causal DAG (directed acyclic graph) for a problem under consideration, is a major roadblock when doing Judea Pearl's Causal Inference (CI) in Statistics. The same problem arises when doing CI in Artificial Intelligence (AI) and Machine Learning (ML). As with many problems in Science, we think Nature has found an effective solution to this problem. We argue that human and animal brains contain an explicit engine for doing CI, and that such an engine uses as input an atlas (i.e., collection) of causal DAGs. We propose a simple algorithm for constructing such an atlas from a library of books or videos/movies. We illustrate our method by applying it to a database of randomly generated Tic-Tac-Toe games. The software used to generate this Tic-Tac-Toe example is open source and available at [GitHub](#).

“If humans were so good at causal inference, religion would not exist.”

- Yann LeCun, *Ref.[1]*

“How much of human knowledge is captured in all the text ever written? To which my answer is: not much. ”

- Yann LeCun, *Ref.[2]*

1 Introduction

Automated Causal DAG Extraction From Text (DEFT) seems like a daunting task today, at a time when little work has been done to achieve it. Some previous work to achieve DEFT can be found in Ref.[7, 4]. Once DEFT is achieved, the extracted DAGs can be collected together as an atlas (i.e., collection) of DAGs. This DAG atlas can be used as input to a Causal Inference (CI) engine within AI/ML software. The job of such a CI engine will be to distinguish between correlation and causation, and to perform all 3 rungs (prediction, do operators, counterfactuals) of Judea Pearl’s ladder of CI. (See Ref. [3]).

Contrary to what naysayers like Yann LeCun predict, we believe an explicit CI engine will some day be an essential part of all AI software.

We also believe that, contrary to what Yann LeCun claims in the above quote, there is ample evidence that a very efficient CI engine already exists in all human and animal brains. Maybe not all human brains excel at CI, but some clearly have (e.g., Einstein’s and Feynman’s). And if human/animal brains possess an explicit CI engine, it’s not too far fetched to assume that they also possess a DAG atlas to provide input to that CI engine. There even seems to be evidence that humans/animals are not born as a Tabula Rasa (blank slate), but possess at birth a primitive DAG atlas. For example, one can find on YouTube videos showing how human babies are born with an aversion to sitting on grass.

Some people, like Yann LeCun, believe that human/animal brains, rather than possessing an explicit CI engine, perform CI through spontaneous generation, or random chance, or Divine Intervention. We vehemently disagree. An AI without an explicit CI engine would be DIM (Divine Intervention Mechanism) witted, and would not act like the human or animal mind. If at the dawn of life, human and animal brains had lacked an explicit CI engine, Nature, through evolution, would have invented it early on. The survival advantage of such a CI engine is enormous.

Finally, contrary to what LeCun claims in the above quote, the knowledge captured in all the text ever written, is enormous. One of the goals of this paper is to improve AI by mining the vast trove of human knowledge (especially, causal knowledge) contained in text (and other sources like videos). Once that vast trove of human knowledge can be mined in an automated fashion, this will open the floodgates

for implementing CI in AI/ML and Statistics. As Judea Pearl has said, CI will elevate AI from Deep Learning to Deep Understanding.

We claim that the DEFT algorithm proposed in this paper extracts *causal* DAGs. This begs the question, what is causality? This is how I define causality in my book “Bayesuivius” (Ref.[5])

Causality is a time-induced ordering between two events, the transmission of information (and its accompanying energy) from the earlier of the two events to the later one, and the physical response of the later event to the reception of that information.

I expand more on this definition in Bayesuivius, in its Appendix E entitled: “Bayesian Networks, Causality and the Passage of Time”.

2 Proposal for a simple DEFT algorithm

In this section, we will present our proposal for a simple DEFT algorithm. We will use the acronym DEFT to mean extraction of DAGs from various primary sources. By **primary sources**, we will mean, not just text (e.g., libraries of books), but also libraries of videos and other sources or a mixture of these.

The first step in our algorithm is to covert the primary sources into a library of one or more comic books (CBs). By a **comic book (CB)**, we mean a book that consists of a finite number of chronologically ordered **frames**. Different CBs in the library may contain a different number of frames. Each frame in a CB must be annotated with an **event descriptor** for each of the most important **events** occurring in the frame. An example of an event descriptor might be “Boy kicks soccer ball”. It’s just a very simple sentence with at least a subject and a verb, and, occasionally, more info than a subject-verb.

Reduction of primary sources to a library of CBs might be done completely by machine learning software (fully-automated) or completely by humans or both. There already exist various image-to-text and text-to-image conversion software apps. If the primary sources are a library of movies, perhaps the image-to-text conversion apps can be modified to select a finite number of chronological ordered frames and to annotate them with event descriptors. If the primary sources are a library of books, perhaps the text-to-image apps can be modified to produce the CBs. Hence, it’s possible that someday a fully automated DEFT algorithm will be achieved.

The reason why we do primary sources to CBs conversion as the first step of our DEFT algorithm is that a primary source like a text often has the chronological order of events implicit. To mine the causal info of a primary source, we first need to make its chronological order of events very clear and explicit. There is much more to causality than chronological ordering, but chronological ordering is an intrinsic component of it.

The second step in our DEFT algorithm is to convert the library of CBs produced in the first step into a collection of DAGs (**DAG atlas**). We shall describe this second step using a simple fictional example.

For our fictional example, we shall end up considering a **number of events** $N_e = 11$. The event descriptors are as follows. Note that sometimes, to save space, we will abbreviate the event descriptor by a 4 letter symbol.

1. *SuRi*: Sun rises
2. *RoCr*: Rooster crows
3. *BuSt*: Bus starts daily rounds
4. *BoWa*: Boy wakes up
5. *BoEa*: Boy eats breakfast
6. *BoRu*: Boy runs to meet bus
7. *BuAr*: Bus arrives at boy's house
8. *DoBa*: Dog barks
9. *BoBo*: Boy boards bus
10. *RaSt*: Rain starts
11. *RaEn*: Rain ends

Suppose two CBs c_1 and c_2 have event sets $E(c_1)$ and $E(c_2)$, respectively. We say c_1 **is time compatible (tc) smaller than** c_2 ($c_1 < c_2$) if (1) $E(c_1)$ is a proper subset of $E(c_2)$ and (2) the events in c_1 are in the same chronological order as the corresponding events in c_2 .

We define a **time compatible CB library (cbLib*)** to be a library (i.e., set) of CBs wherein there is one CB of maximum length, and all other CBs in the library are tc-smaller than the maximum length one.

A CB c **is time compatible with a cbLib*** C if $c < c_{max}$, where c_{max} is the maximum length CB of C .¹

Given a large set of CBs, we can generate a collection \mathcal{C} of cbLib*s as follows. If \mathcal{C} is empty, find a CB c , then create a cbLib* with c as its only CB, and add that cbLib* to \mathcal{C} . If \mathcal{C} is not empty, and we want to add a new CB called c to \mathcal{C} , add c to every cbLib* in \mathcal{C} with which it is time compatible. If we can't find a cbLib* in \mathcal{C} with which c is time compatible, then create a new cbLib* with c as its only member and add it to \mathcal{C} .

¹Another possible definition for this is that $c < c_{max}$ or $c > c_{max}$.

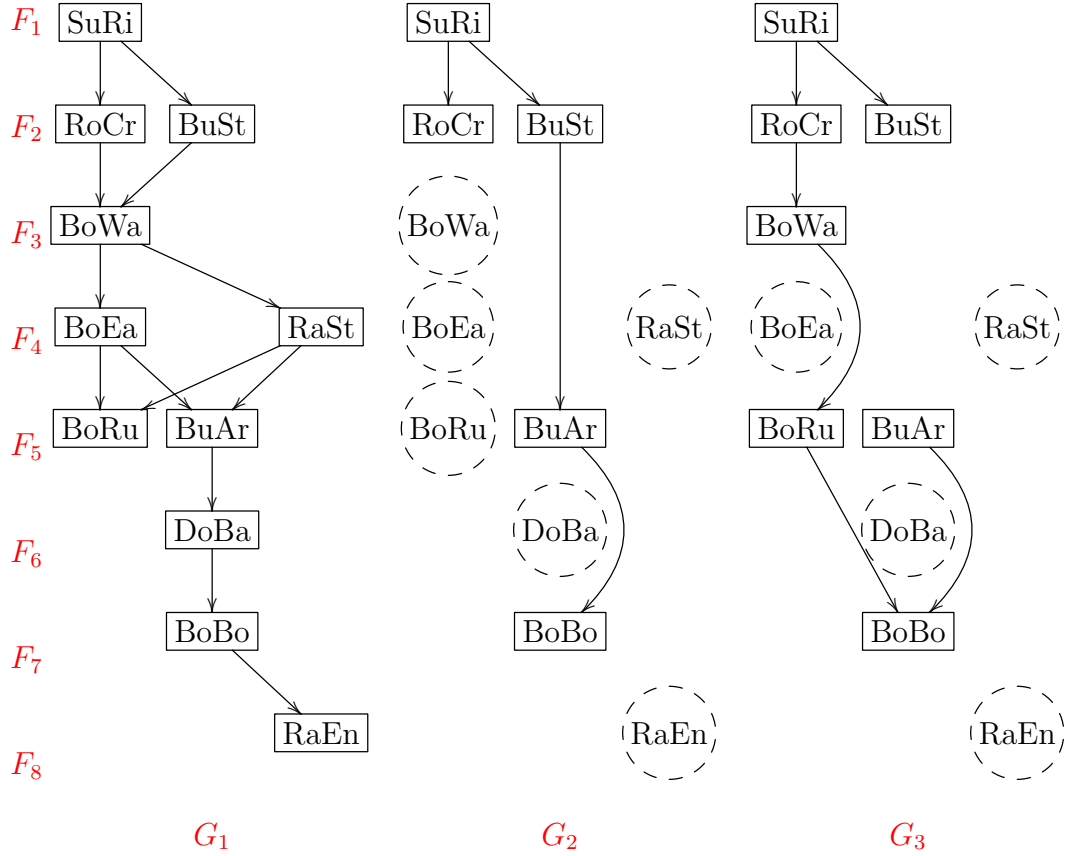


Figure 1: Let $\gamma = 1, 2, 3, \dots, N_{cb}$ label the CBs in our current $cbLib^*$. For the γ 'th CB, we construct a DAG G_γ . Here we show the DAGs for the first 3 CBs of our current $cbLib^*$. Each row F_i for $i = 1, 2, \dots, N_f$ are the frames of the CBs. Nodes enclosed by a rectangle represent events that happened in a CB. Nodes enclosed by a dashed circle are placeholders for events that didn't happen in a CB.

Next we shall explain how to process each $cbLib^*$ individually.

Let $\gamma = 1, 2, 3, \dots, N_{cb}$ label the CBs in our current $cbLib^*$. For the γ 'th CB, we construct a DAG G_γ . Fig.1 shows, for our fictional example, the DAGs for the first 3 CBs of our current $cbLib^*$. The DAGs in Fig.1 were constructed very easily by placing the event descriptors in chronological order with time pointing down the page, and adding arrows entering each event e from all events e' in the frame immediately before e 's frame. We will refer to this as $T_{mem} = 1$ **memory time**. We could have just as well added arrows entering each event e from all events e' in the two frames immediately before e 's frame. This would be the case $T_{mem} = 2$. In general, one could use any T_{mem} , including T_{mem} as big as the total number of frames, but the higher T_{mem} is, the more dense the DAG will be. I think T_{mem} equal to 1 or 2 will

provide a good enough causal fit most of the time.

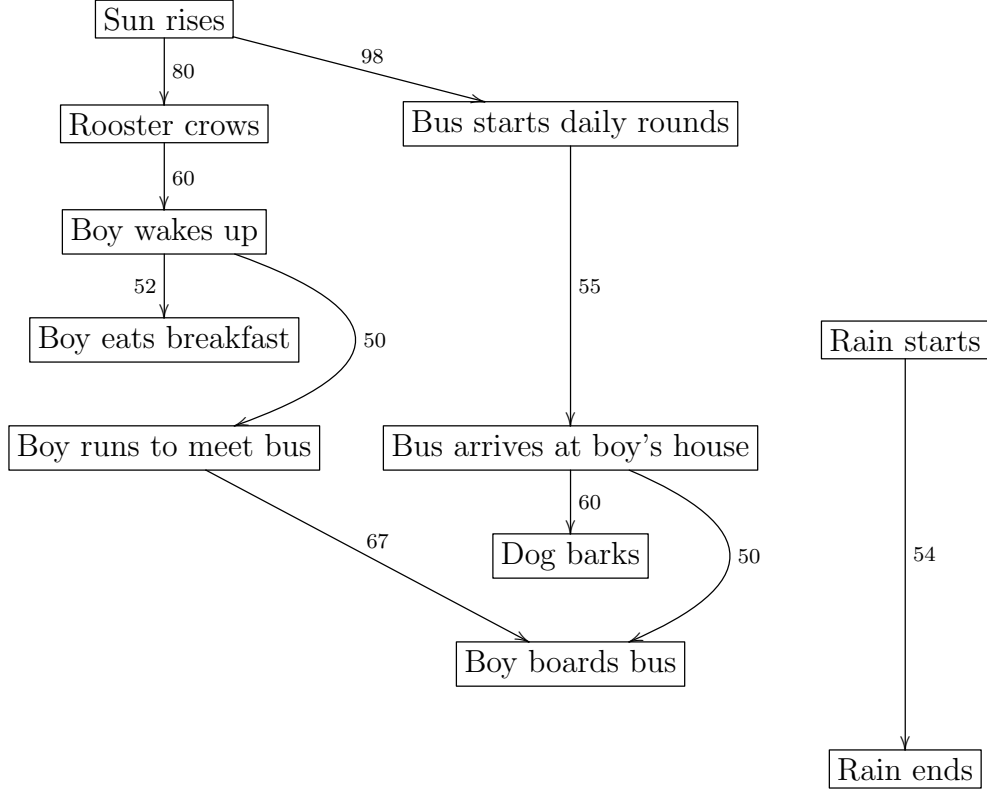


Figure 2: High frequency arrows DAG, G_{hfa} , with $N_{art} = 50$. Arrow repetition number given besides each arrow. Fig.1 shows DAGs G_γ for $\gamma = 1, 2, 3$.

At this point, we are ready to calculate G_{hfa} , the **high frequency arrows (hfa)** DAG. This DAG is created by counting the number of times each arrow appears in the just created DAGs G_γ for all γ , and keeping only arrows that are repeated at least as much as the **arrow repetition threshold** N_{art} . Fig.2 shows a possible G_{hfa} for our fictional example, assuming $N_{art} = 50$. The intuition behind G_{hfa} is that arrows that have low repetition numbers do so because Nature is telling us that they are not essential for the causal mechanism being expressed by the whole DAG.

Note that one can upgrade G_{hfa} from a causal DAG to a causal Bayesian Network (bnet) by calculating an empirical estimate for the TPM (Transition Probability Matrix) of each event node $\underline{e} \in \{0, 1\}$ of G_{hfa} . This estimate can be calculated using the following formula.

$$P(e|pa(\underline{e}) = \pi) = \frac{\sum_{\gamma=1}^{N_{cb}} \mathbb{1}(\underline{e}_\gamma = e, [pa(\underline{e})]_\gamma = \pi)}{\sum_{\gamma=1}^{N_{cb}} \mathbb{1}([pa(\underline{e})]_\gamma = \pi)} \quad (1)$$

Here $e \in \{0, 1\}$ and $\pi \in \{0, 1\}^n$, where n is the number of parents of node \underline{e} of G_{hfa} .

Also, node \underline{e} in G_{hfa} and node \underline{e}_γ in G_γ refer to the same event.²

Note that G_{hfa} might consist of more than one connected DAG. Furthermore, note that we have shown how to generate a G_{hfa} for each $cbLib^*$ in a collection \mathcal{C} of $cbLib^*$ s. Hence, this algorithm generates a collection of G_{hfa} DAGs. This collection of DAGs is what we call the **DAG atlas**.

Note that our DEFT algorithm is very simple and only requires counting, and keeping good chronologically ordered records. We believe that the human brain excels at those tasks. Our DEFT algorithm avoids performing precise numerical calculations like calculating π to 15 decimal places because we believe most human brains can't perform such tasks.

3 How to use the DAG Atlas

Once a DAG atlas is calculated, and its DAGs are upgraded to bnets, it can be used to perform predictions (1st rung), do operators (2nd rung), and to calculate CATE and ATE for counterfactuals (3rd rung). The results of these calculations can be used by the AI to make decisions.

In the particular version of the 3 rungs that is advocated in Bayesuivius³, the 2nd rung removes arrows from a bnet using a do-operator, and the 3rd rung adds new nodes and arrows to a bnet using an imagine operator. My point is that the DAG atlas can be modified by applying the rules of CI.

Note that hidden nodes have not been mentioned so far. Recall that hidden nodes are nodes in a DAG that one would like to include because we think they are important, but which haven't been measured, either because it is impossible to measure them, or we simply haven't gotten around to measuring them. Hidden nodes could certainly be added to the DAGs of a DAG atlas, either in an automated fashion, or by hand.

Note that if a CB repeats a task over and over again, the DAG G_{hfa} might turn out to be a dynamical Bayesian network.⁴

A Appendix: Tic-Tac-Toe example

See Ref.[6] for Python code that implements, for Tic-Tac-Toe, the DEFT algorithm proposed in this paper.

A CB is a list of chronologically ordered frames. For the Tic-Tac-Toe example considered in this code, a CB is one Tic-Tac-Toe game, and a frame is just one move. For example,

²We are using the convention used in Bayesuivius (Ref.[5]) of indicating random variables by underlining them, and using $pa(\underline{x})$ to denote the parent nodes of node \underline{x} .

³See the chapter entitled "Counterfactual Reasoning" in Bayesuivius, Ref. [5]

⁴Dynamical Bayesian networks are discussed in Bayesuivius, Ref.[5]

$$['X1', 'O2', 'X8', 'O0', 'X6', 'O4', 'X7'] \quad (2)$$

represents a CB/game and 'X1' is a frame/move. The X and O refer to the player. X always plays first. The numbers refer to the positions on the Tic-Tac-Toe grid, labelled as follows:

$$\begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array} \quad (3)$$

In the more general case considered in this paper, a frame, instead of being a single string, can be a list of strings called event descriptors, which represent simultaneous events in the frame.

The time each X or O is played will be indicated by the time enclosed in parenthesis. For example, X3(2) is an X played at time 2 and position 3. The 3, which indicates the position on the Tic-Tac-Toe grid, is redundant and omitted if X3(2) is rendered inside a Tic-Tac-Toe grid. For example, here is the game Eq.(2) represented on a grid.

$$\begin{array}{|c|c|c|} \hline O(3) & X(0) & O(1) \\ \hline & O(5) & \\ \hline X(4) & X(6) & X(2) \\ \hline \end{array} \quad (4)$$

Our code generates 2000 CBs at random, and calculates a collection \mathcal{C} of cbLib*s from those. For each cbLib* $C \in \mathcal{C}$, it can draw a DAG for each CB $c \in C$, for a given memory time T_{mem} . It can then calculate the high frequency arrows DAG G_{hfa} of C , for a given arrow repetition threshold N_{art} . It can also upgrade G_{hfa} to a bnet.

References

- [1] Yann LeCun. Twitter, all world text quote. <https://twitter.com/ylecun/status/1562137291845521408>.
- [2] Yann LeCun. Twitter, causal inference and religion quote. <https://twitter.com/ylecun/status/1577128801620070400>.
- [3] Judea Pearl and Dana Mackenzie. *The book of why: the new science of cause and effect*. Basic books, 2018.
- [4] John L. Pfaltz. Logical implication and causal dependency. in Conceptual structures: inspiration and application, Springer Verlag LNAI, volume 4068, pages 145–157, year 2006, <https://www.cs.virginia.edu/~jlp/06.ICCS.pdf>. I found out about Pfaltz's paper on Twitter, from Boris Sobolev, after the present paper was finished.

- [5] Robert R. Tucci. Bayesuvius (book). <https://github.com/rrtucci/Bayesuvius/raw/master/main.pdf>.
- [6] Robert R. Tucci. `deft-tic-tac-toe` at github. <https://github.com/rrtucci/deft-tic-tac-toe>.
- [7] Justin Wood, Nicholas Matiasz, Alcino Silva, William Hsu, Alexej Abyzov, and Wei Wang. Opberg: Discovering causal sentences using optimal alignments. In *International Conference on Big Data Analytics and Knowledge Discovery*, pages 17–30. Springer, 2022.